

验证 MySQL MGR 双机房双活架构可行性

原创 洪斌 玩转MySQL 2025年05月06日 11:23 四川

在实际生产环境中，很多企业都希望实现**双机房双活**，以提升业务的高可用性和容灾能力。恰好最近某客户问到：**MGR 如何实现双机房双活？**

MySQL Group Replication (MGR) 作为官方的高可用方案，支持多主模式、异步复制链路故障转移 (Asynchronous Replication Channel Failover)。那么，在双活场景下是否可行？有哪些坑？

为此，我用cursor写了个测试脚本，来快速验证 MGR 架构下双机房双活的可行性。

设计目标

- 验证 MySQL MGR 在双机房部署下，能否实现双活（即两地同时可写、数据自动同步、故障自动切换）。
- 自动化部署、切换、故障模拟、数据一致性校验 减少人工干预，提升测试效率

主要功能

1. 一键部署双集群

- 支持单主 / 多主两种模式
- 使用 **dbdeployer** 快速搭建两套 MGR 集群，模拟双机房环境

2. 灵活复制链路测试

- 可配置单向或双向复制
- 自动建立异步复制通道，模拟双活场景下的数据写入

3. 自动化故障与恢复

- 支持主节点故障模拟
- 自动等待新主选举
- 节点重启与重新加入集群
- 完整还原生产故障场景

4. 数据一致性与复制状态校验

- 自动插入测试数据
- 校验两地数据同步情况
- 实时监控复制链路健康状态

测试结论

通过该脚本的自动化测试，得出如下结论：

- **单主模式**下不建议配置双向复制，推荐使用 **MySQL Shell 部署 ClusterSet 架构**，通过 **MySQL Router 实现两集群双活**
- 单主或多主模式下若配置双向复制，**必须启用** **skip_replica_start**，否则节点重启后在加入集群前会先同步对方集群的数据，导致与本组 GTID 不一致，无法加入集群
- **MySQL Shell 暂不支持配置双向复制**，只能部署 InnoDB ClusterSet 架构

使用方法示例

```
### 部署单主模式集群
python mgr_test.py -s

### 部署多主模式集群
python mgr_test.py -m

### 运行单主单向复制测试
python mgr_test.py -ss

### 运行多主双向复制测试
python mgr_test.py -aa
```

脚本源码

```
1  #!/usr/bin/env python3
2
3  from typing import List, Dict, Literal
4  import subprocess
5  import time
6  from dataclasses import dataclass
7  from datetime import datetime
8  import argparse
9  MYSQL_VERSION = "8.0.40"
10 REPL_USER = "msandbox"
11 REPL_PASSWORD = "msandbox"
12 # MGR集群端口配置
13 MGR1_PORTS = [4316, 4317, 4318] # 第一个MGR集群的端口
14 MGR2_PORTS = [5316, 5317, 5318] # 第二个MGR集群的端口
15 SANDBOX= "$HOME/workbench/sandboxes"
16 REPLICATION_CHANNELS = {
```

```
17     "mgr1": "async_mgr2_to_mgr1", # mgr1 集群上的通道
18     "mgr2": "async_mgr1_to_mgr2"  # mgr2 集群上的通道
19 }
20 @dataclass
21 class MGRNode:
22     port: int
23     host: str = '127.0.0.1'
24     user: str = 'msandbox'
25     password: str = 'msandbox'
26
27     def execute_sql(self, sql: str, vertical: bool = False) -> str:
28         cmd = f"mysql -h{self.host} -P{self.port} -u{self.user} -p{self.password} {sql}"
29         if not vertical:
30             cmd += " -NB"
31         cmd += f" -e \"{sql}\""
32         try:
33             result = subprocess.run(cmd, shell=True, capture_output=True, text=True)
34             return result.stdout.strip()
35         except subprocess.CalledProcessError as e:
36             error_msg = e.stderr.strip()
37             print("\n❌ 最近的错误日志:")
38             self.show_error_log()
39             raise Exception(e.stderr)
40
41     def show_error_log(self, lines: int = 20):
42         """显示最近的错误日志"""
43
44         Args:
45             lines: 显示的日志行数
46
47         """
48         try:
49             # 获取错误日志文件路径
50             log_file = f"{SANDBOX}/mgr{1 if self.port < 5000 else 2}/error.log"
51             cmd = f"tail -n {lines} {log_file} | grep ERROR"
52             result = subprocess.run(cmd, shell=True, capture_output=True, text=True)
53             print(result.stdout.strip())
54         except subprocess.CalledProcessError as e:
55             print(f"⚠️ 无法读取错误日志: {e.stderr.strip()}")
56
57     def shutdown(self, force: bool = False):
58         if force:
59             # 只杀死MySQL进程，而不是所有监听该端口的进程
60             cmd = f"ps aux | grep mysql.*{self.port} | grep -v grep | awk '{print $2}' | xargs kill -9"
61         else:
62             cmd = f"mysqladmin -u{self.user} -p{self.password} -h{self.host} shutdown"
63         subprocess.run(cmd, shell=True, check=True)
```

```

63 class MGRCluster:
64     def __init__(self, name: str, ports: List[int], mysql_version: str):
65         self.name = name
66         self.nodes = [MGRNode(port) for port in ports]
67         self.primary_node = self.nodes[0]
68         self.mysql_version = mysql_version
69         self._primary_node_cache = None
70         self._last_primary_check = 0
71         self._cache_ttl = 10 # 缓存10秒
72
73     def deploy(self, topology_mode: Literal["single-primary", "multi-
74         # 先删除已存在的集群
75         try:
76             print(f"=== 删除集群 {self.name} ===")
77             cmd = f"ps aux | grep {self.name} | grep -v grep | awk '{
78             cmd = f"dbdeployer delete {self.name} --concurrent 2>/dev
79             subprocess.run(cmd, shell=True, check=False)
80         except:
81             pass # 忽略删除失败
82
83         # 部署新的MGR集群
84         print(f"=== 部署集群 {self.name} ===")
85         cmd = f"dbdeployer deploy replication --topology=group -c ski
86         if topology_mode == "single-primary":
87             cmd += " --single-primary"
88         cmd += f" --nodes={len(self.nodes)} --concurrent --port-as-se
89         print(f"{cmd}")
90         subprocess.run(cmd, shell=True, check=True)
91         # 清除缓存
92         self._primary_node_cache = None
93         self._last_primary_check = 0
94     def get_primary_node(self, force_check: bool = False) -> MGRNode:
95         """获取主节点
96
97         Args:
98             force_check: 是否强制检查，忽略缓存
99         """
100         current_time = time.time()
101
102         # 如果缓存有效且不强制检查，直接返回缓存的主节点
103         if not force_check and self._primary_node_cache and \
104             (current_time - self._last_primary_check) < self._cache_tt
105             return self._primary_node_cache
106         print(f"\n🔍 开始查找集群 {self.name} 的主节点")
107
108         # 先检查组复制状态

```

```

109         for node in self.nodes:
110             try:
111                 print(f" ⚡ 检查节点 {node.port}")
112                 time.sleep(3)
113                 group_status = node.execute_sql("""
114                     SELECT COUNT(*) FROM performance_schema.replicati
115                     WHERE MEMBER_STATE = 'ONLINE'
116                 """)
117                 print(f" 📊 在线节点数: {group_status}")
118
119                 # 如果有节点在线，从这个节点查询主节点信息
120                 if int(group_status) > 0:
121                     members = node.execute_sql("""
122                         SELECT CONCAT(MEMBER_PORT, ' (', MEMBER_STATE
123                         FROM performance_schema.replication_group_mem
124                     """)
125                     print(f" 📋 节点状态: {members.replace('\n', ', ')}")
126
127                     # 查找主节点
128                     for check_node in self.nodes:
129                         role = check_node.execute_sql("""
130                             SELECT MEMBER_ROLE
131                             FROM performance_schema.replication_group
132                             WHERE MEMBER_ID = @@server_uuid
133                         """)
134                         if role == "PRIMARY":
135                             self.primary_node = check_node
136                             # 更新缓存
137                             self._primary_node_cache = check_node
138                             self._last_primary_check = current_time
139                             print(f" ✅ 找到主节点: {check_node.port}")
140                             return check_node
141                     break # 如果找到在线节点但没找到主节点，不需要继续检查其他
142             except subprocess.CalledProcessError as e:
143                 print(f" ❌ 检查节点 {node.port} 失败")
144                 continue
145
146         # 清除缓存
147         self._primary_node_cache = None
148         self._last_primary_check = 0
149         raise Exception(f"! 集群 {self.name} 中未找到主节点，可能正在进行主")
150
151     def init_test_schema(self):
152         """初始化测试数据库和表"""
153         primary = self.get_primary_node()
154         primary.execute_sql("""

```

```

155         CREATE DATABASE IF NOT EXISTS mgr_test;
156         CREATE TABLE IF NOT EXISTS mgr_test.events (
157             id INT AUTO_INCREMENT PRIMARY KEY,
158             event_time DATETIME,
159             cluster_name VARCHAR(10),
160             event_type VARCHAR(20),
161             event_data VARCHAR(100)
162         );
163     """
164     def switch_mode(self, to_single_primary: bool = True):
165         """切换MGR模式"""
166         primary = self.get_primary_node()
167         mode = "单主" if to_single_primary else "多主"
168         print(f"\n🔄 切换{self.name}到{mode}模式")
169
170         try:
171             mode_sql = "SINGLE_PRIMARY" if to_single_primary else "MULTI_MASTER"
172             primary.execute_sql(f"""
173                 SELECT group_replication_switch_to_{mode_sql.lower()}.
174             """)
175             # 清除缓存强制重新检查主节点
176             self._primary_node_cache = None
177             self._last_primary_check = 0
178             print(f"✅ {self.name}已切换到{mode}模式")
179         except Exception as e:
180             print(f"❌ 切换失败: {str(e)}")
181             raise
182
183     class ReplicationManager:
184     def __init__(self, source_cluster: MGRCluster, target_cluster: MGRCluster):
185         self.source = source_cluster
186         self.target = target_cluster
187
188     def setup_replication(self, source: MGRNode, target: MGRNode, channel_name: str):
189         channel_clause = f"FOR CHANNEL '{channel_name}'" if channel_name else ""
190         auto_failover_clause = ",\n                SOURCE_CONNECTION_AUTO_FAIL_OVER=1"
191         target.execute_sql(f"""
192             CHANGE REPLICATION SOURCE TO
193             SOURCE_HOST='{source.host}',
194             SOURCE_PORT={source.port},
195             SOURCE_USER='{REPL_USER}',
196             SOURCE_PASSWORD='{REPL_PASSWORD}',
197             SOURCE_CONNECT_RETRY=3,
198             SOURCE_AUTO_POSITION=1{auto_failover_clause}
199             FOR CHANNEL '{channel_name}';
200             START REPLICA {channel_clause};
201         """)

```

```
201
202     def setup_async_replication(self):
203         source = self.source.get_primary_node()
204         target = self.target.get_primary_node()
205         # 使用全局变量中的通道名称
206         self.setup_replication(source, target, REPLICATION_CHANNELS[s
207
208     def setup_active_active_replication(self):
209         source_primary = self.source.get_primary_node()
210         target_primary = self.target.get_primary_node()
211
212         # 使用全局变量中的通道名称
213         self.setup_replication(source_primary, target_primary,
214                                REPLICATION_CHANNELS[self.target.name],
215                                auto_failover=False)
216         self.setup_replication(target_primary, source_primary,
217                                REPLICATION_CHANNELS[self.source.name],
218                                auto_failover=False)
219 # 部署相关功能
220 class ClusterDeployer:
221     def __init__(self, mysql_version: str = MYSQL_VERSION):
222         self.mysql_version = mysql_version
223
224     def deploy_single_primary_clusters(self, cluster1_ports: List[int]
225         """部署两个单主模式MGR集群"""
226         cluster1 = MGRCluster("mgr1", cluster1_ports, self.mysql_vers
227         cluster2 = MGRCluster("mgr2", cluster2_ports, self.mysql_vers
228
229         print("\n🚀 开始部署测试集群")
230         cluster1.deploy("single-primary")
231         cluster2.deploy("single-primary")
232
233         print("\n🗄️ 初始化测试数据库")
234         cluster1.init_test_schema()
235         cluster2.init_test_schema()
236
237         return cluster1, cluster2
238
239     def deploy_multi_primary_clusters(self, cluster1_ports: List[int]
240         """部署两个多主模式MGR集群"""
241         cluster1 = MGRCluster("mgr1", cluster1_ports, self.mysql_vers
242         cluster2 = MGRCluster("mgr2", cluster2_ports, self.mysql_vers
243
244         print("\n🚀 开始部署多主模式测试集群")
245         cluster1.deploy("multi-primary")
246         cluster2.deploy("multi-primary")
```

```

247
248     print("\n📁 初始化测试数据库")
249     cluster1.init_test_schema()
250     cluster2.init_test_schema()
251
252     return cluster1, cluster2
253 # 测试相关功能
254 class ClusterTester:
255     def __init__(self, cluster1: MGRCluster = None, cluster2: MGRCluster = None):
256         self.cluster1 = cluster1
257         self.cluster2 = cluster2
258
259     def _check_clusters(self):
260         """检查集群是否已部署"""
261         try:
262             # 检查 dbdeployer 中的沙箱列表
263             cmd = "dbdeployer sandboxes --header | grep -E 'mgr[12]'"
264             result = subprocess.run(cmd, shell=True, capture_output=True)
265             if result.stdout.strip():
266                 # 找到已部署的集群，初始化集群对象
267                 self.cluster1 = MGRCluster("mgr1", MGR1_PORTS)
268                 self.cluster2 = MGRCluster("mgr2", MGR2_PORTS)
269                 return
270             except subprocess.CalledProcessError:
271                 pass
272
273             raise Exception("❌ 未找到已部署的集群，请先使用 's' 或 'm' 命令部署")
274
275     @staticmethod
276     def insert_test_data(node: MGRNode, cluster_name: str, event_type: str, data: str):
277         """插入测试数据
278
279         Args:
280             node: 目标节点
281             cluster_name: 集群名称
282             event_type: 事件类型（如：初始数据、故障后数据等）
283             data: 事件数据
284         """
285         node.execute_sql(f"""
286             INSERT INTO mgr_test.events (event_time, cluster_name, event_type, data)
287             VALUES (NOW(), '{cluster_name}', '{event_type}', '{data}')
288         """)
289
290     @staticmethod
291     def verify_data_sync(node: MGRNode):
292         """验证数据同步状态"""

```



```

293         result = node.execute_sql("""
294             SELECT CONCAT(
295                 '时间: ', event_time,
296                 ', 集群: ', cluster_name,
297                 ', 类型: ', event_type,
298                 ', 数据: ', event_data
299             ) FROM mgr_test.events ORDER BY event_time
300         """)
301         print(f"\n📊 在节点 {node.port} 的数据: \n{result}")
302
303     @staticmethod
304     def verify_replication_status(node: MGRNode, channel: str = ""):
305         """验证复制状态"""
306         print(f"\n🔍 验证复制通道 '{channel}' 状态")
307         cmd = f"mysql -h{node.host} -P{node.port} -u{node.user} -p{no
308         try:
309             result = subprocess.run(cmd, shell=True, capture_output=T
310             print(result.stdout)
311         except subprocess.CalledProcessError as e:
312             print(f"❌ 检查复制状态失败: {e.stderr if e.stderr else '通道
313             raise
314
315     def setup_replication(self, mode: Literal["async", "active-active
316         """设置复制关系"""
317         self._check_clusters()
318         repl_mgr = ReplicationManager(self.cluster1, self.cluster2)
319
320         mode_name = "单向" if mode == "async" else "双向"
321         print(f"\n🔄 设置{mode_name}复制")
322         if mode == "async":
323             repl_mgr.setup_async_replication()
324             self.verify_replication_status(self.cluster2.get_primary_
325                                     REPLICATION_CHANNELS[self.cl
326         else:
327             repl_mgr.setup_active_active_replication()
328             for cluster in [self.cluster1, self.cluster2]:
329                 self.verify_replication_status(cluster.get_primary_no
330                                     REPLICATION_CHANNELS[clu
331     def write_test_data(self, cluster: MGRCluster, event_type: str =
332         """写入测试数据"""
333         primary = cluster.get_primary_node()
334         print(f"\n📝 在{cluster.name}的主节点{primary.port}写入测试数据")
335         self.insert_test_data(primary, cluster.name, event_type,
336                               f"来自节点{primary.port}")
337         time.sleep(1) # 间隔写入以区分时间顺序
338     def simulate_node_failure(self, node: MGRNode, cluster: MGRCluste

```

```

339         """模拟节点故障
340
341     Args:
342         node: 要故障的节点
343         cluster: 节点所属集群
344
345     Returns:
346         MGRNode: 新的主节点（如果发生切换）
347     """
348     print(f"\n✳ 模拟故障：关闭节点 {node.port}")
349     node.shutdown(force=True)
350
351     # 如果故障节点是主节点，等待新主选举
352     if node.port == cluster.primary_node.port:
353         return self._wait_for_new_primary(cluster)
354     return cluster.get_primary_node()
355 def recover_node(self, node: MGRNode, cluster: MGRCluster):
356     """恢复故障节点
357
358     Args:
359         node: 要恢复的节点
360         cluster: 节点所属集群
361     """
362     print(f"\n🔄 重启节点 {node.port} 并重新加入集群")
363     node_index = cluster.nodes.index(node) + 1
364     subprocess.run(f"{SANDBOX}/{cluster.name}/node{node_index}/start",
365                   shell=True, check=True)
366     time.sleep(2) # 等待启动完成
367     print(f"\n🔄 节点: {node.port} START GROUP_REPLICATION 重新加入集群")
368     node.execute_sql("START GROUP_REPLICATION;")
369     print(f"\n🔄 节点: {node.port} START REPLICA 开启异步复制")
370     node.execute_sql("START REPLICA;")
371     time.sleep(10) # 等待节点加入
372 def verify_cluster_data(self, clusters: List[MGRCluster]):
373     """验证集群数据同步状态"""
374     print("\n🔍 验证数据同步")
375     time.sleep(5) # 等待数据同步
376
377     for cluster in clusters:
378         primary = cluster.get_primary_node()
379         # 验证数据同步
380         self.verify_data_sync(primary)
381
382         # 验证复制状态 - 使用固定的通道名称
383         self.verify_replication_status(primary, REPLICATION_CHANNEL)
384 def test_async_failover(self):

```

```
385         """测试单向复制故障转移"""
386         # 设置单向复制
387         self.setup_replication("async")
388
389         # 写入初始数据
390         self.write_test_data(self.cluster1, "初始数据")
391         time.sleep(2) # 等待复制
392
393         # 模拟主节点故障
394         primary1 = self.cluster1.get_primary_node()
395         new_primary = self.simulate_node_failure(primary1, self.cluster1)
396
397         # 在新主写入数据
398         self.write_test_data(self.cluster1, "故障转移后")
399         time.sleep(2)
400
401         # 恢复故障节点
402         self.recover_node(primary1, self.cluster1)
403
404         # 验证数据同步
405         self.verify_cluster_data([self.cluster1, self.cluster2])
406     def test_active_active_replication(self):
407         """测试双向复制"""
408         # 设置双向复制
409         self.setup_replication("active-active")
410
411         # 在两个集群写入数据
412         for cluster in [self.cluster1, self.cluster2]:
413             self.write_test_data(cluster, "初始数据")
414
415         # 模拟节点故障和恢复
416         primary1 = self.cluster1.get_primary_node()
417         new_primary = self.simulate_node_failure(primary1, self.cluster1)
418
419         # 继续写入数据
420         self.write_test_data(self.cluster1, "故障之后")
421         self.write_test_data(self.cluster2, "故障期间")
422
423         # 恢复故障节点
424         self.recover_node(primary1, self.cluster1)
425
426         # 验证数据同步
427         self.verify_cluster_data([self.cluster1, self.cluster2])
428     def _wait_for_new_primary(self, cluster: MGRCluster, max_wait: int):
429         """等待新主节点选举完成"""
430         print("\n⌚ 等待新主节点选举")
```

```

431     new_primary = None
432     old_primary_port = cluster.primary_node.port
433
434     for i in range(max_wait // interval):
435         try:
436             print(f"\n🔄 第 {i + 1} 次尝试查找 {cluster.name} 新主节点,
437             # 强制检查主节点, 忽略缓存
438             new_primary = cluster.get_primary_node(force_check=True)
439             if new_primary and new_primary.port != old_primary_port:
440                 print(f"✅ 新主节点选举成功: {new_primary.port}")
441                 cluster.primary_node = new_primary # 更新集群的主节点
442                 return new_primary
443         except Exception as e:
444             # print(f"⌚ 等待主节点选举: \n{str(e)}")
445             time.sleep(interval)
446             continue
447
448         raise Exception(f"❌ 在 {max_wait} 秒内未完成主节点选举")
449 def deploy_clusters(topology: Literal["single-primary", "multi-primary"],
450                    """部署MGR集群"""
451                    deployer = ClusterDeployer()
452                    if topology == "single-primary":
453                        return deployer.deploy_single_primary_clusters(MGR1_PORTS, MGR2_PORTS)
454                    else:
455                        return deployer.deploy_multi_primary_clusters(MGR1_PORTS, MGR2_PORTS)
456 def run_tests(test_type: Literal["async", "active-active"] = "active-active",
457              """运行指定类型的测试"""
458              tester = ClusterTester(*clusters) if clusters else ClusterTester(*clusters)
459              if test_type == "async":
460                  tester.test_async_failover()
461              else:
462                  tester.test_active_active_replication()
463 if __name__ == "__main__":
464     parser = argparse.ArgumentParser(description="MGR集群部署和测试工具")
465     parser.add_argument('-s', '--deploy-single', action='store_true',
466     parser.add_argument('-m', '--deploy-multi', action='store_true',
467     parser.add_argument('-ss', '--single-async', action='store_true',
468     parser.add_argument('-sa', '--single-active', action='store_true',
469     parser.add_argument('-aa', '--multi-active', action='store_true',
470
471     args = parser.parse_args()
472
473     # 如果没有提供任何参数, 显示帮助信息
474     if not any(vars(args).values()):
475         parser.print_help()
476         exit(0)

```

```
477
478 # 映射参数到场景
479 if args.deploy_single:
480     topology, test_type = "single-primary", None
481 elif args.deploy_multi:
482     topology, test_type = "multi-primary", None
483 elif args.single_async:
484     topology, test_type = "single-primary", "async"
485 elif args.single_active:
486     topology, test_type = "single-primary", "active-active"
487 elif args.multi_active:
488     topology, test_type = "multi-primary", "active-active"
489
490 clusters = None
491 tester = None
492
493 # 如果是测试场景，先检查已部署的集群
494 if test_type:
495     try:
496         tester = ClusterTester()
497         tester._check_clusters()
498         clusters = (tester.cluster1, tester.cluster2)
499
500         # 根据测试场景切换集群模式
501         to_single = topology == "single-primary"
502         current_mode = "单主" if to_single else "多主"
503         print(f"\n=== 切换到{current_mode}模式 ===")
504         for cluster in clusters:
505             cluster.switch_mode(to_single)
506     except Exception as e:
507         print(f"\n⚠️ 未找到已部署的集群或切换模式失败，重新部署集群")
508         clusters = None
509
510 # 如果是部署命令或者没有找到已部署的集群，部署新集群
511 if not clusters:
512     print(f"\n=== 部署{topology}模式集群 ===")
513     clusters = deploy_clusters(topology)
514
515 # 如果有测试类型，运行测试
516 if test_type:
517     print(f"\n=== 运行测试场景: {test_type} ===")
518     run_tests(test_type, clusters)
```

