

MySQL生产实战优化（利用Index skip scan优化性能提升257倍）

原创 陈刚 春田花花程序园 2025年03月31日 15:12 四川

背景

我们生产环境均采用单表多租户设计模式，即一个表中存有多个租户数据，在实际的业务中每个租户的数据是隔离的，即A租户不可能访问到B租户数据，假设租户ID的字段为profileid，我们每个索引都会将profileid放在最左第一列(主键除外)，如（profileid，date,status）。

昨天运营需要统计一个线上数据，会跨租户查询,where条件中不能带有租户ID字段，这将导致所有索引均失效，整个SQL查询非常非常慢。

问题SQL

实际语句更复杂，我将语句简单化了，说明优化过程

研发同学写的原版SQL

```
explain
select  profileid,billcode,billtype,billdate,billid from a1
where a1.billtype IN( 604) and a1.billdate >='2025-03-14' and  a1.billdate
<='2025-03-15'
```

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	a1		ALL					3113812	1.11	Using where

通过执行计划看到，将会是全表扫描，扫描300多万行

实际执行需要46327ms

```
--> Filters: ((a1.billtype = 604) and (a1.billdate >= TIMESTAMP'2025-03-14 00:00:00') and (a1.billdate <= TIMESTAMP'2025-03-15 00:00:00')) (cost=568326.61 rows=34581) (actual time=46006.572..46327.106 rows=502 loops=1)
--> Table scan on a1 (cost=568326.61 rows=3112920) (actual time=0.102..45871.070 rows=3344248 loops=1)
```

性能优化

优化思路：

实际上此表有（profileid,billtype,billdate）联合索引，而billid 是主键字段，在整个语句中只有billcode不在上述联合索引之内，所以改造成先走Index skip scan，然后再自关联取billcode即可。

一个表中租户一般低于1000,所以完全符合Index skip scan的条件。我来改写，让这个语句走Index skip scan

优化的第一版SQL

```
explain
select
    t.*,bi.billcode from
(
```

```

select  profileid,billtype,billdate,billid from a1
where a1.billtype IN( 604) and a1.billdate >='2025-03-14' and a1.billdate
<='2025-03-15'
) t
inner join erp_bill_index a1 as bi on t.billid = bi.billid

```

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	a1		index	PRIMARY,idx_year	ix_profiled_billtype_billdate	13		1113873	1.11	Using index conditions; Using index
1	SIMPLE	bi		eq_ref	PRIMARY,idx_year	PRIMARY	8	erp_bill_index.a1.billid	1	100.00	

通过执行计划发现，MySQL给我合并了子查询，虽然扫描行数未变，但走了索引覆盖扫描
虽然没有走到 index skip scan 但性能还是有较大提升 从46327ms提升到了546ms 提升84倍

```

--> Nested loop inner join (cost=913444.46 rows=3113458) (actual time=93.717..549.420 rows=530 loops=1)
--> Index scan on a1 using ix_profiled_billtype_billdate, with index conditions: ([a1.billtype = 604] and [a1.billdate >= TIMESTAMP'2025-03-14 00:00:00'] and [a1.billdate <= TIMESTAMP'2025-03-15 00:00:00'] (cost=568302.27 rows=3113458) (actual time=93.469..546.476 rows=530 loops=1)
--> Single-row index lookup on bi using PRIMARY (billid=a1.billid) (cost=0.99 rows=1) (actual time=0.005..0.005 rows=1 loops=530)

```

优化的第二版SQL

```

explain
select /*+no_merge() */
t.*,bi.billcode from
(
select  profileid,billtype,billdate,billid from a1
where a1.billtype IN( 604) and a1.billdate >='2025-03-14' and a1.billdate
<='2025-03-15'
) t
inner join a1 bi on t.billid = bi.billid

```

强制用hint no_merge() 阻止MySQL的合并优化

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	PRIMARY	<derived2>		ALL					57655	100.00	
1	PRIMARY	bi		eq_ref	PRIMARY,idx_year	PRIMARY	8	t.billid	1	100.00	
2	DERIVED	a1		range	ix_profiled_billtype_billdate...	ix_profiled_billtype_billdate	8		519004	11.11	Using where; Using index for skip scan

在执行计划中看到扫描行数少了，也终于走到了 Index skip scan

性能继续提升 从46327ms提升到了178ms 提升260倍

```

--> Nested loop inner join (actual time=174.886..178.174 rows=536 loops=1)
--> Table scan on t (actual time=0.001..0.096 rows=536 loops=1)
--> Materialize (actual time=174.849..175.041 rows=536 loops=1)
--> Filter: ([a1.billtype = 604] and [a1.billdate >= TIMESTAMP'2025-03-14 00:00:00'] and [a1.billdate <= TIMESTAMP'2025-03-15 00:00:00']) (cost=133529.30 rows=57647) (actual time=31.086..174.658 rows=536 loops=1)
--> Index range scan on a1 using index_for_skip_scan(ix_profiled_billtype_billdate) (cost=133529.30 rows=518929) (actual time=0.031..132.366 rows=287581 loops=1)
--> Single-row index lookup on bi using PRIMARY (billid=t.billid) (cost=0.97 rows=1) (actual time=0.005..0.005 rows=1 loops=536)

```

Index skip scan的要求

- 1、索引为复合索引
- 2、查询为单表查询
- 3、查询不使用GROUP BY或DISTINCT
- 4、查询仅引用索引中的列
- 5、索引为 (a,b,c) 条件用到了b,c 且b、c设计不能为null

贴上官方的说明

Using this strategy decreases the number of accessed rows because MySQL skips the rows that do not qualify for each constructed range. This Skip Scan access method is

applicable under the following conditions:

Table T has at least one compound index with key parts of the form ([A_1, ..., A_k,] B_1, ..., B_m, C [, D_1, ..., D_n]). Key parts A and D may be empty, but B and C must be nonempty.

The query references only one table.

The query does not use GROUP BY or DISTINCT.

The query references only columns in the index.

The predicates on A_1, ..., A_k must be equality predicates and they must be constants.

This includes the IN() operator.

The query must be a conjunctive query; that is, an AND of OR conditions:

(cond1(key_part1) OR cond2(key_part1)) AND (cond1(key_part2) OR ...) AND ...

There must be a range condition on C.

Conditions on D columns are permitted. Conditions on D must be in conjunction with the range condition on C.

Use of Skip Scan is indicated in EXPLAIN output as follows:

Using index for skip scan in the Extra column indicates that the loose index Skip Scan access method is used.

If the index can be used for Skip Scan, the index should be visible in the possible_keys column.

range-access-skip-scan